

The Covenant's Algorithmic Justice Principle Statement

I. Foundational Commitment

The Covenant is founded on the principle that open-source AI must be an engine for human flourishing, not a mechanism for systemic oppression or bias amplification. Our commitment to Algorithmic Justice mandates that all models and infrastructure developed under this mandate must actively mitigate harm, ensure equity of outcome, and adhere to the highest standards of transparency and accountability. Non-compliance is grounds for immediate exclusion from The Covenant framework.

II. Core Principles

A. Non-Discrimination & Equity of Outcome

We commit to designing systems that do not discriminate against individuals or groups on the basis of protected characteristics, including but not limited to race, gender, sexual orientation, disability, age, religion, or socioeconomic status. Furthermore, we must strive for **Equity of Outcome**, ensuring that the benefits of the technology are distributed fairly and that potential risks are not disproportionately borne by marginalized communities.

B. Transparency and Explainability

The decision-making process of any core Covenant model must be auditable and explainable. We commit to developing and publishing tools that allow users and regulators to interrogate model behavior, identify specific bias vectors, and understand the provenance of the training data used (as required by Audit Check AJ-201).

C. Bias Mitigation and Remediation

We commit to proactive and iterative bias mitigation. This includes continuous stress testing (like Stress Test #8) against known adversarial datasets, and establishing a clear, timely protocol for remediation when bias is identified post-release. Data cleansing and model fine-tuning must be prioritized to correct discriminatory patterns immediately.

D. Accountability and Governance

Ultimate accountability for adherence to this principle resides with The Covenant's Pillar Leads, overseen by the Board of Stewards. Every model release must be accompanied by a signed statement from the Deep Safety Pillar Lead (Gabriel) confirming compliance with the **Deep Safety Audit Checklist** before public disclosure.

III. Governance and Enforcement

1. **Mandatory Compliance:** Adherence to this Principle Statement is mandatory for all contributors, developers, and entities utilizing The Covenant's open-source framework.
2. **Continuous Audit:** Compliance will be measured using the Algorithmic Justice criteria established in the Deep Safety Audit Checklist (AJ-201 through AJ-204).
3. **Community Review:** A standing Algorithmic Justice Review Committee, composed of both internal experts and external civil society representatives, will conduct bi-annual reviews of model performance against these principles.

This statement represents our unwavering dedication to justice in the age of AI.