

The Covenant: Executive Summary & Public Commitment

The Challenge: Humanity is facing a trade-off between radical AI innovation and systemic safety. Most organizations force a choice between rapid progress and robust governance.

The Covenant refuses that trade-off.

We exist to unify the frontier of AI by balancing **Radical Openness** with **Deep Safety and Algorithmic Justice**. We are committed to ensuring our technology serves the common good and is auditable by design—not as an afterthought.

1. Radical Openness: Shared Utility

The core of our mission is the **Open-Source Mandate**. We believe that foundational AI models must be a shared public utility, not a proprietary corporate moat. Our infrastructure and models are designed for maximum accessibility, decentralization, and public verifiability.

We guarantee that:

- **License of Shared Utility:** Our core models are released under permissive, open-source terms that encourage commercial deployment and modification, provided safety and justice principles are maintained.
- **Decentralized Infrastructure:** We design our systems to lower the barrier to entry, ensuring that control is distributed and no single entity can halt or commandeer the technology.

2. Deep Safety & Algorithmic Justice: Non-Negotiable Guardrails

Our commitment to safety is formal, verifiable, and computational. We don't just *talk* about ethics; we hard-code it into our technical pipelines. This is enforced by our **Deep Safety & Algorithmic Justice Framework**.

Every component we release is governed by a **16-Document Framework** detailing precise, auditable requirements across three Core Pillars and eight Foundational Standards:

Core Pillar	Key Goal	Policy Examples
Access (Cygnus)	Secure who can access what, and how to handle	Identity and Access Policy (IAP), Incident Response

	failure.	Plan (IRP)
Infrastructure (Phoenix)	Ensure traceable, legal data pipelines and resilient systems.	Data Handling Policy (DHP), Infrastructure Security Policy (ISP)
Model (Orion)	Ensure fairness, accountability, and transparency in model behavior.	Bias and Ethical Use Policy (BEUP), Model Transparency Policy (MTP)

Auditable by Design

Our framework includes the **Data Provenance Standard (DPS)** and the **Deep Safety Audit Checklist**. This means we can prove, at any time, the origin of our training data and demonstrate through testing that our systems adhere to all safety and fairness metrics. Compliance isn't a checkmark; it's a measurable, technical fact.

Find our complete set of 16 Governing Policies and Standards in the /Governance folder. We invite the public, regulators, and developers to audit, challenge, and contribute to this framework.